

- Problem 3

- What are your observations about the query plans for the scanning and sorting of such differently sized bags S?
 - Below is time to scan the entire data, the scan time grows in relation to size of S i.e $O(N)$. The planning time seems constant.

SCAN				SORT			
Data Output	Explain	Messages	Notifications	Scratch Pad	Data Output	Explain	Messages
size of relation S integer	Avg execution time to scan S (in ms): SELECT x FROM S;	numeric			size of relation S integer	Avg execution time to sort S (in ms): SELECT x FROM S order by	numeric
1	10	0.007			1	10	0.019
2	100	0.012			2	100	0.032
3	1000	0.068			3	1000	0.194
4	10000	0.658			4	10000	2.490
5	100000	6.579			5	100000	29.459
6	1000000	65.986			6	1000000	343.130
7	10000000	737.383			7	10000000	3854.534
8	100000000	7773.957			8	100000000	30024.235

- What do you observe about the execution time to sort S as a function of n?
 - For relatively small data and based on 'work_mem' parameter, the quick sort is used to sort data but as soon as the data size grows instead of quick sort the postgres uses external merge sort, with time complexity of $O(n \log_B(n))$ where B is block size
- Does this conform with the formal time complexity of (external) sorting? Explain.
 - Since external sort works in time complexity of $O(n \log_B(n))$ where B is block size. Here we take example of block size as 64kb = 64000 bytes thus to sort 100000 records we will take $O(100000 \log_{64000}(100000))$ * time for I/O operation we get corresponding time. All the above experiments are done with work_mem setting of 1.5 GB
- It is possible to set the working memory of PostgreSQL using the set work mem command. For example, set work mem = '16MB' sets the working memory to 16MB.8 The smallest possible working memory in postgresQL is 64kB and the largest depends on your computer memory size. But you can try for example 1GB. Repeat question 3a for memory sizes 64kB and 1GB and report your observations.

64 KB Scan				64 KB Sort			
Data Output	Explain	Messages	Notifications	Scratch Pad	Data Output	Explain	Messages
size of relation S integer	Avg execution time to scan S (in ms): SELECT x FROM S;	numeric			size of relation S integer	Avg execution time to sort S (in ms): SELECT x FROM S order by	numeric
1	10	0.026			1	10	0.015
2	100	0.012			2	100	0.028
3	1000	0.070			3	1000	0.620
4	10000	0.647			4	10000	3.353
5	100000	6.599			5	100000	39.664
6	1000000	66.202			6	1000000	309.161
7	10000000	744.759			7	10000000	3131.792
8	100000000	7869.021			8	100000000	47874.703

1 Gb Scan				1 Gb Sort			
-----------	--	--	--	-----------	--	--	--

Data Output	Explain	Messages	Notifications	Scratch Pad
size of relation S integer	Avg execution time to scan S (in ms): SELECT x FROM S	numeric		
1	10	0.007		
2	100	0.012		
3	1000	0.068		
4	10000	0.675		
5	100000	6.629		
6	1000000	66.294		
7	10000000	745.419		
8	100000000	7896.808		

size of relation S integer	Avg execution time to sort S (in ms): SELECT x FROM S order by	numeric
1	10	0.016
2	100	0.027
3	1000	0.199
4	10000	2.642
5	100000	29.376
6	1000000	345.323
7	10000000	3851.359
8	100000000	40658.351

Since the work_mem parameter sets the block size with which the postgres can work, setting the large block size certainly helps the sorting when data size grows since postgres can accommodate more data in memory but for smaller sizes of data it hardly makes any difference. In case of sorting data since the external merge sort has $\log N$ to base B in time complexity calculation the large size of B reduces the total time for execution.

- I was unable to gather data for sizes above 10^7 for index queries since my system kept crashing. So below is data for index queries.

Create Index

Data Output

Explain

Messages

Notifications

Scratch Pad

size n of relation S

integer

avg execution time to create index indexedS

numeric

1

10

0.802

2

100

1.936

3

1000

14.114

4

10000

143.435

5

100000

1035.816

6

1000000

6969.548

7

10000000

51799.617

Sort Index

Data Output

Explain

Messages

Notifications

Scratch Pad

size n of relation S

integer

avg execution time to create index indexedS

numeric

1

10

7329.073

2

100

7275.081

3

1000

7247.048

4

10000

7289.490

5

100000

7292.106

6

1000000

7262.839

7

10000000

7376.858

- The B+ Tree creation takes time of order $O(N)$ thus we can see exponential growth in index creation.
- The B+ Tree has advantage in terms of range comparison, however with each range search after reading the respective pointers to record we must fetch actual data as well thus we are able to see some overhead in term of range search.

- Problem 7
 - block size = 4096 bytes
 - block-address size = 9 bytes
 - block access time (I/O operation) = 10 ms (micro seconds)
 - record size = 150 bytes
 - record primary key size = 8 bytes

Here if n is number of keys a particular node in B+ tree can hold then we can have

$$8*n + 9*(n + 1) \leq 4096$$



Thus, we can have $n=240$ keys in any B+ tree node.

1. To find any record we will require $\log_n N$ (+ 1 IO operation to read the actual record)


Here the total number of keys is $n = 240$ and total number of records can be 10^9

Thus, total time to find the key will be $\log_{240} 10^9 = 9 * \log_{240} 10 = 3.781179 * \text{block access time} = 3.781179 * 10 = 37.81 \text{ ms}$

2. To insert any record in B+ Tree it will take $\log_n N = \log_{240} 10^9 = 9 * \log_{240} 10 = 3.781179 * \text{block access time} = 3.781179 * 10 = 37.81 \text{ ms}$
3. Memory requirements
 - a. The first level
 - i. The root node will have 240 keys with 241 pointers to child node thus to store root node the memory require will be 4096 bytes
 - b. The second level:
 - i. Since each pointer in root node points to another node with 240 keys and 241 pointers the total memory will be $241 * 4096$ bytes
 - ii. Thus, to store level 1 and level 2 we will require: $4096 + 241 * 4096 = 4096 * 242$ bytes = 0.9453125 MB
 - c. The third level
 - i. At second level we have total of 241 nodes with each node having another 241 blocks pointing to next level thus we have total of
 1. $241 * 241$ blocks at third level
 2. Thus total memory required will be: $241 * 241 * 4096 = 237,899,776$ bytes = 226.87890625 MB.
 3. To store all three levels total memory requirements will be: $0.9453125 + 226.87890625 = 227.82$ MB





INSERT 4




INSERT 10

OVERFLOW






INSERT 12





INSERT 1

OVERFLOW


OVERFLOW







INSERT 7



INSERT 5

OVERFLOW

OVERFLOW

8

INSERT 5

4 7 8

11


1 2


4 5

7


8 10

11 12







INSERT 5




Delete 12




Delete 2




Delete 11




Delete 5



Delete 1



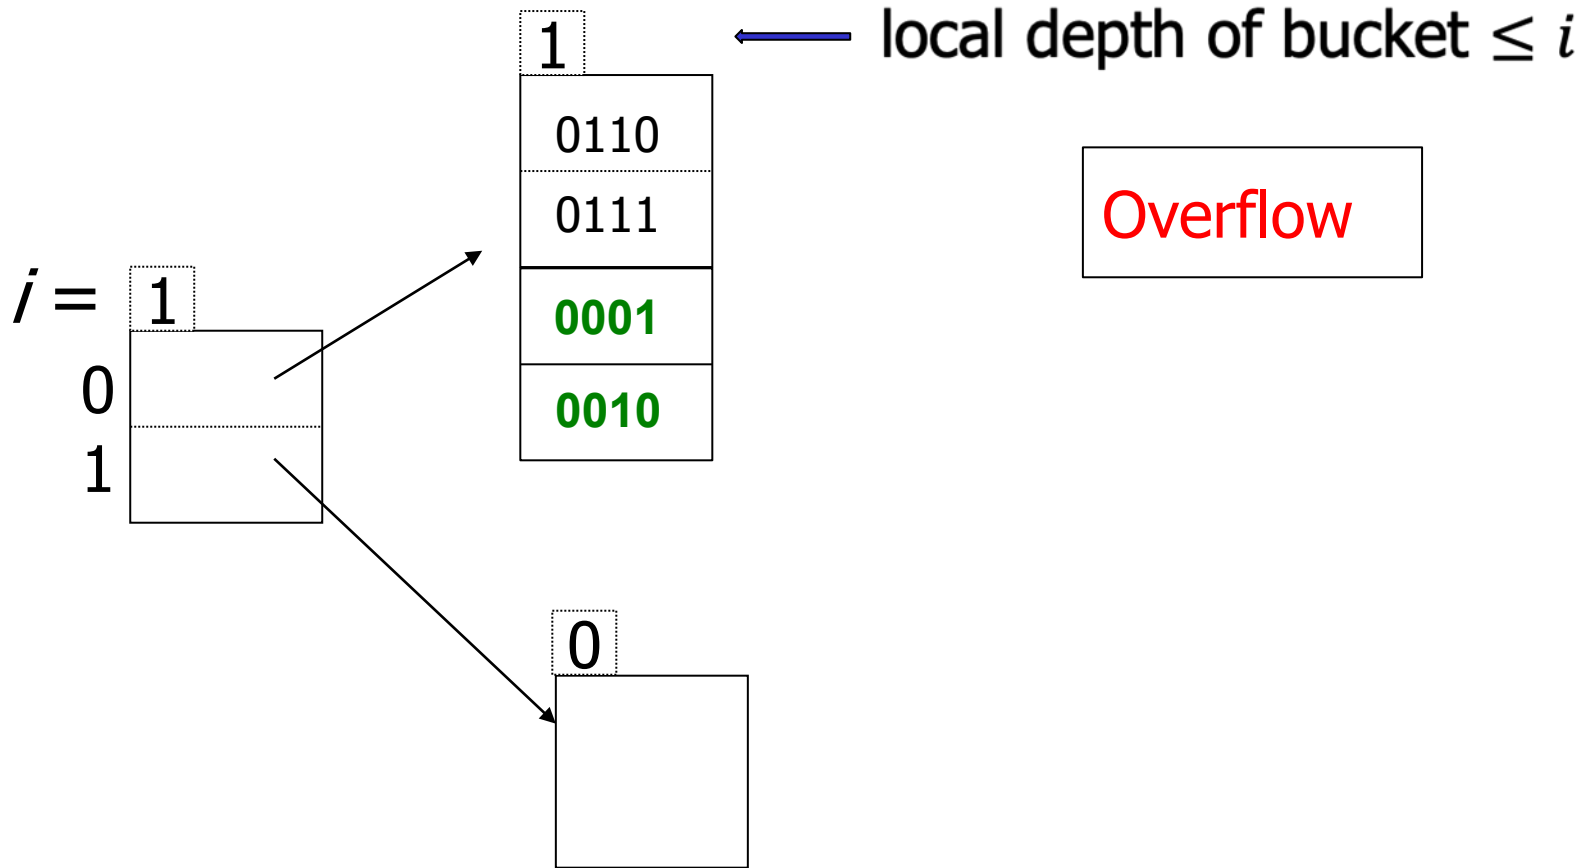
Delete 4



local depth of bucket $\leq i$

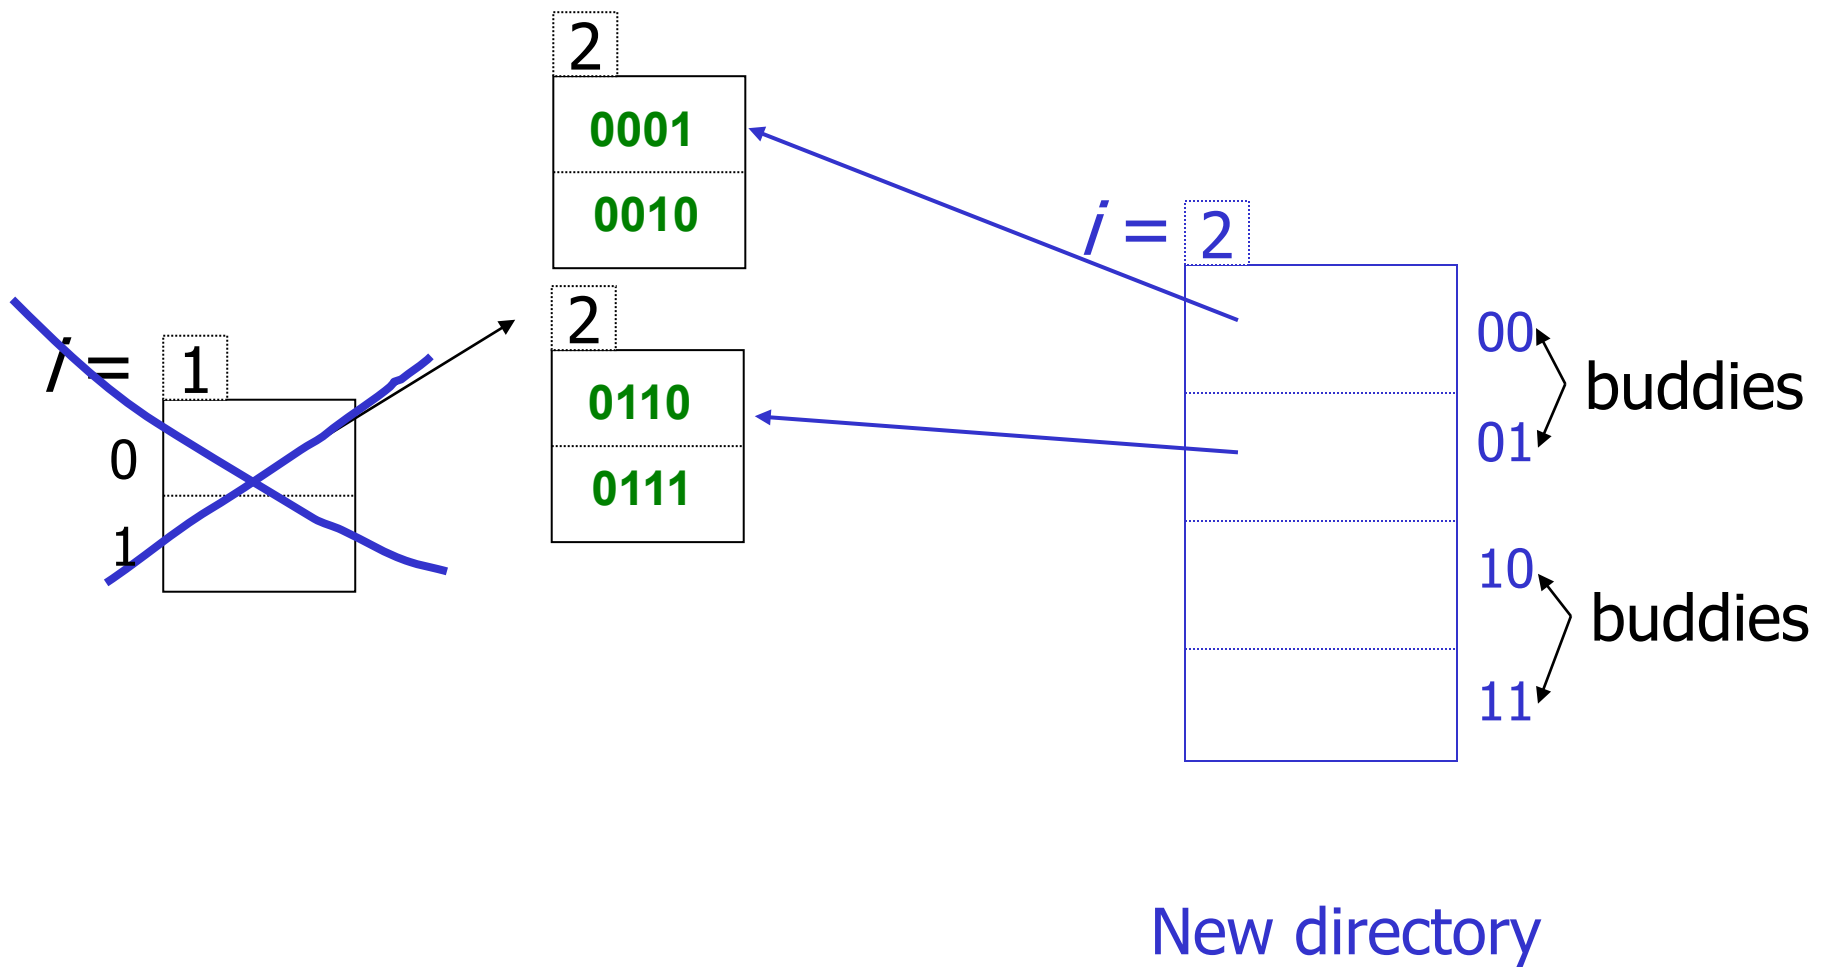
Insert **0110** = 6

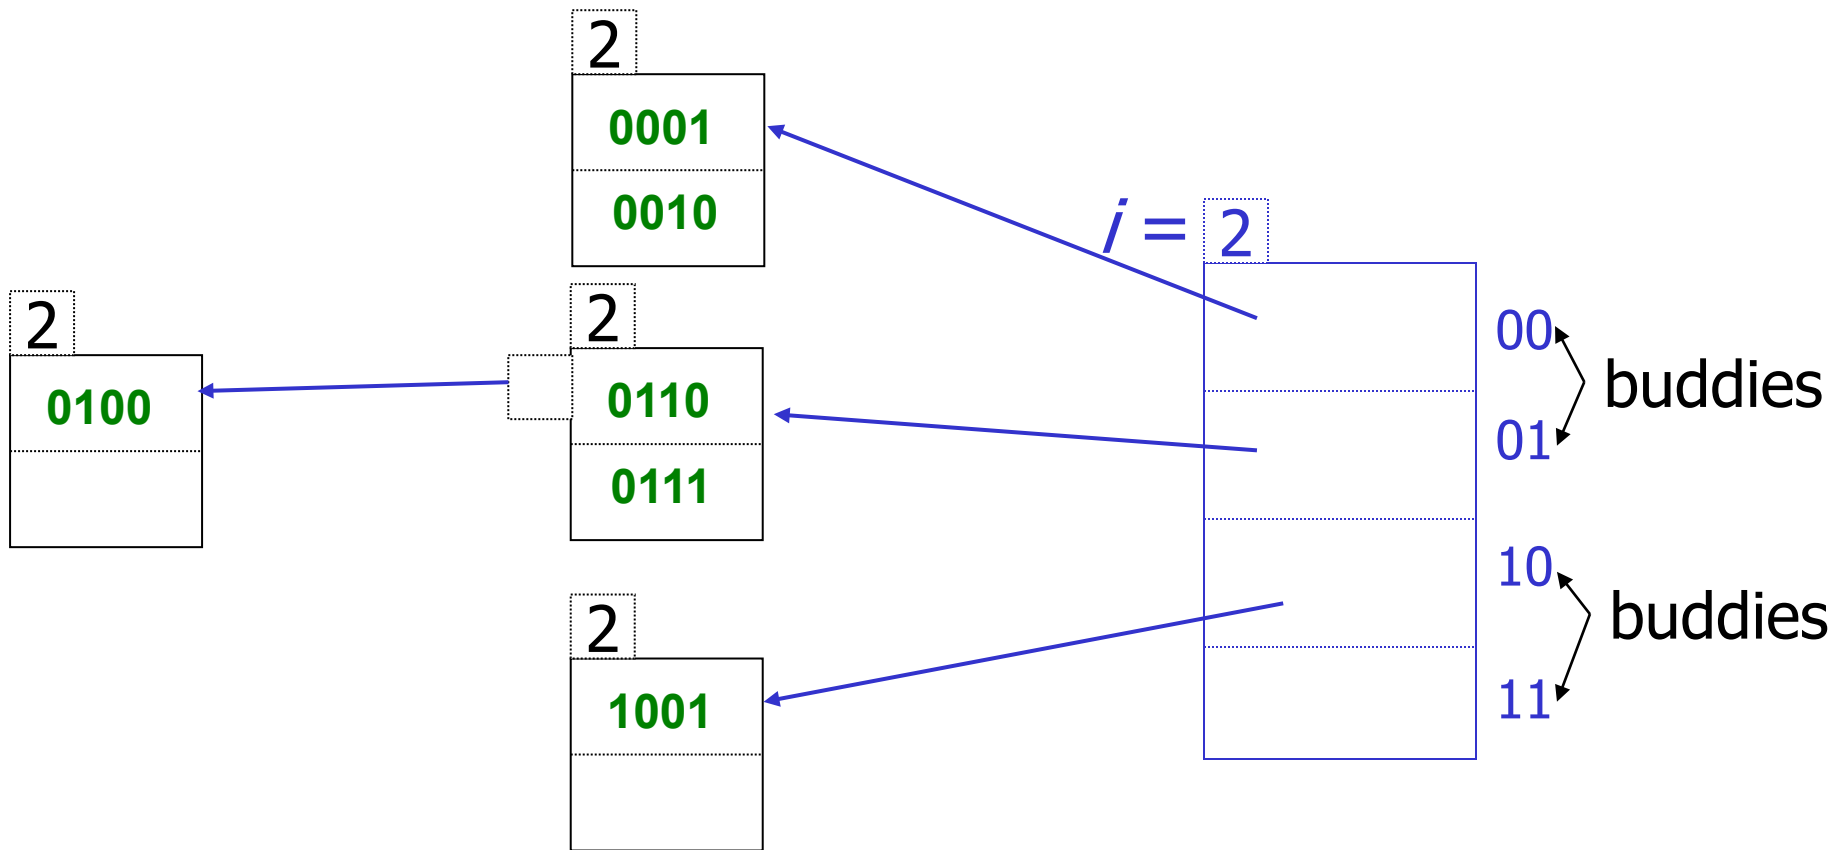
Insert **0111** = 7



Insert 0001 = 1

Insert 0010 = 2

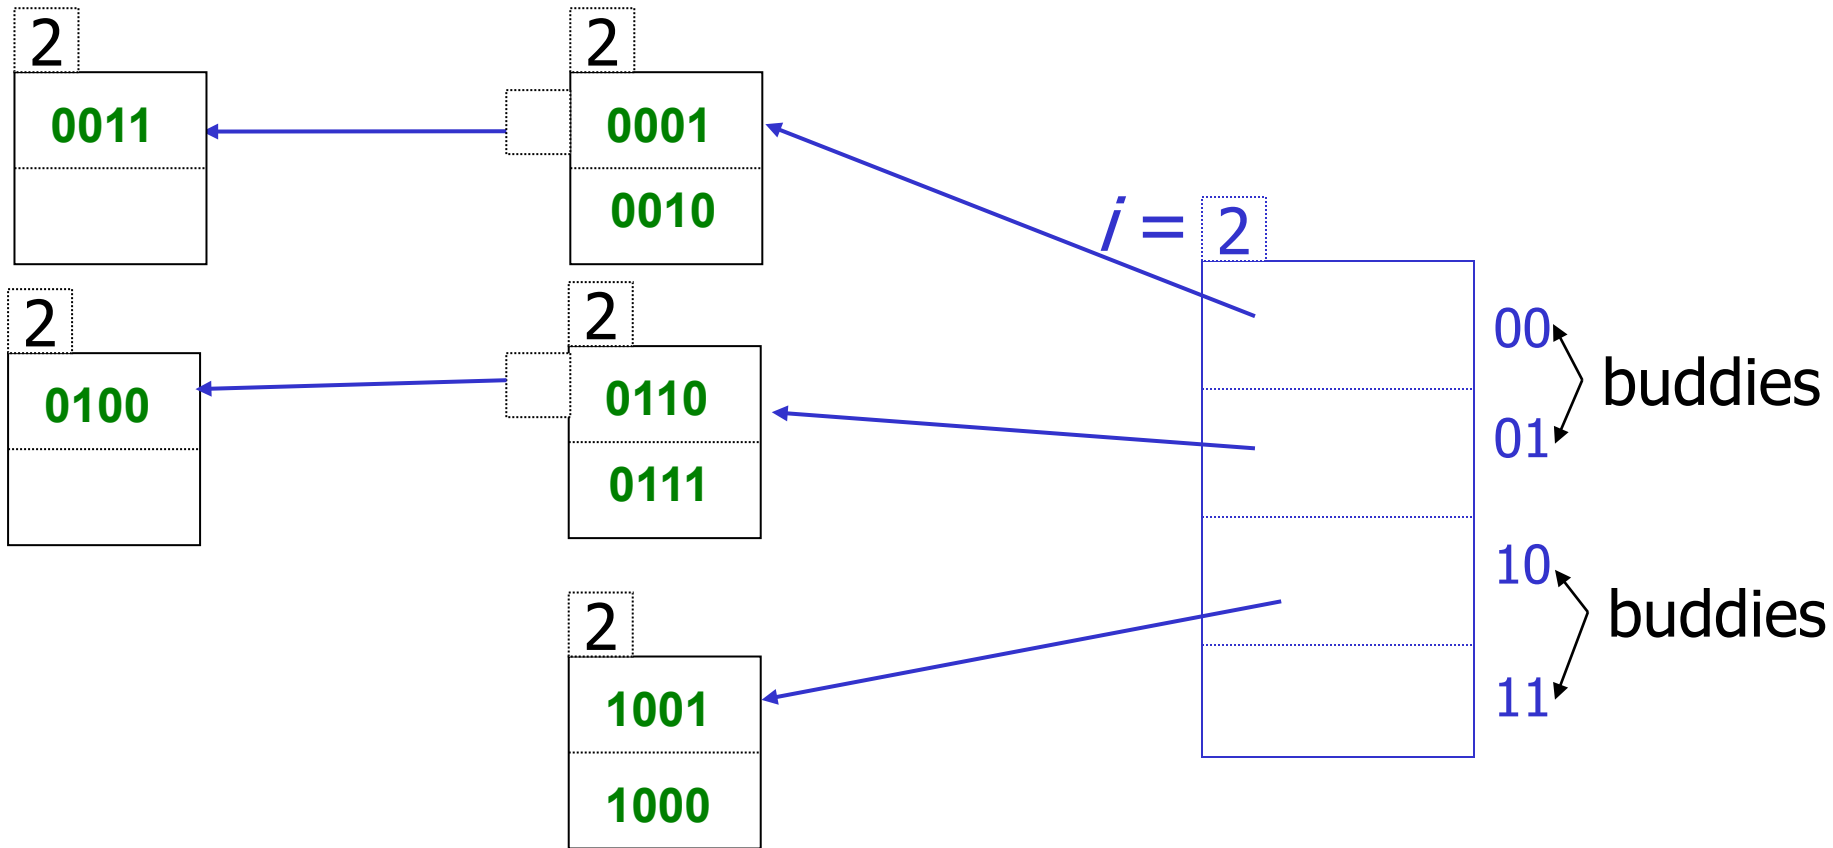




New directory

Insert **1001** = 9

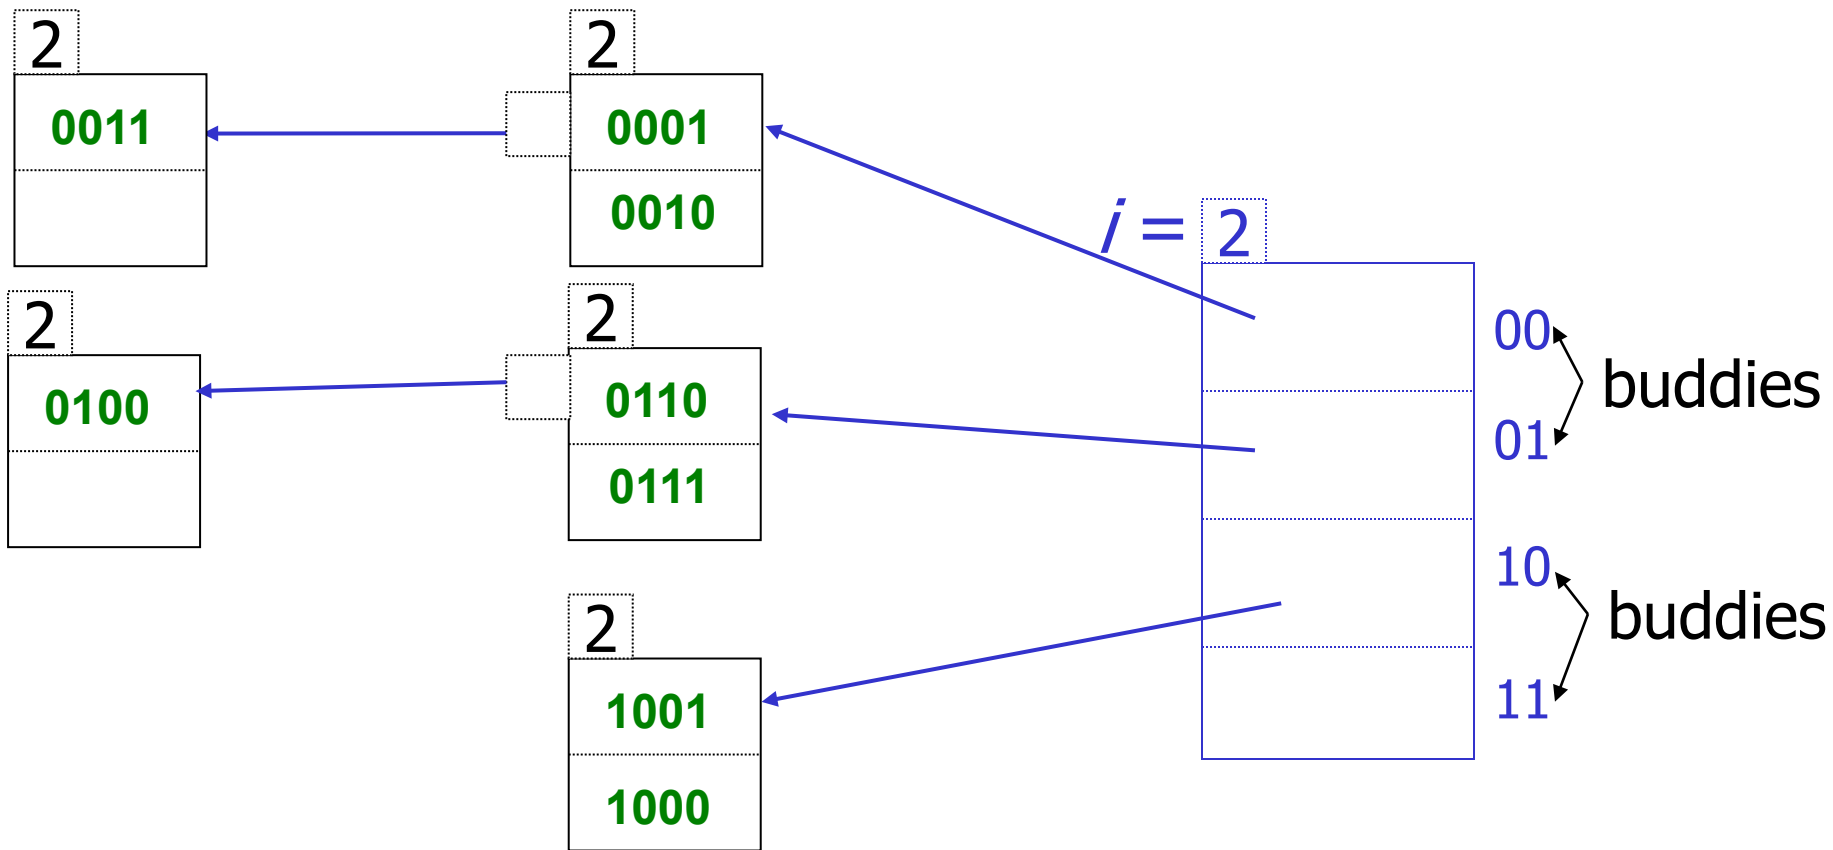
Insert **0100** = 4



New directory

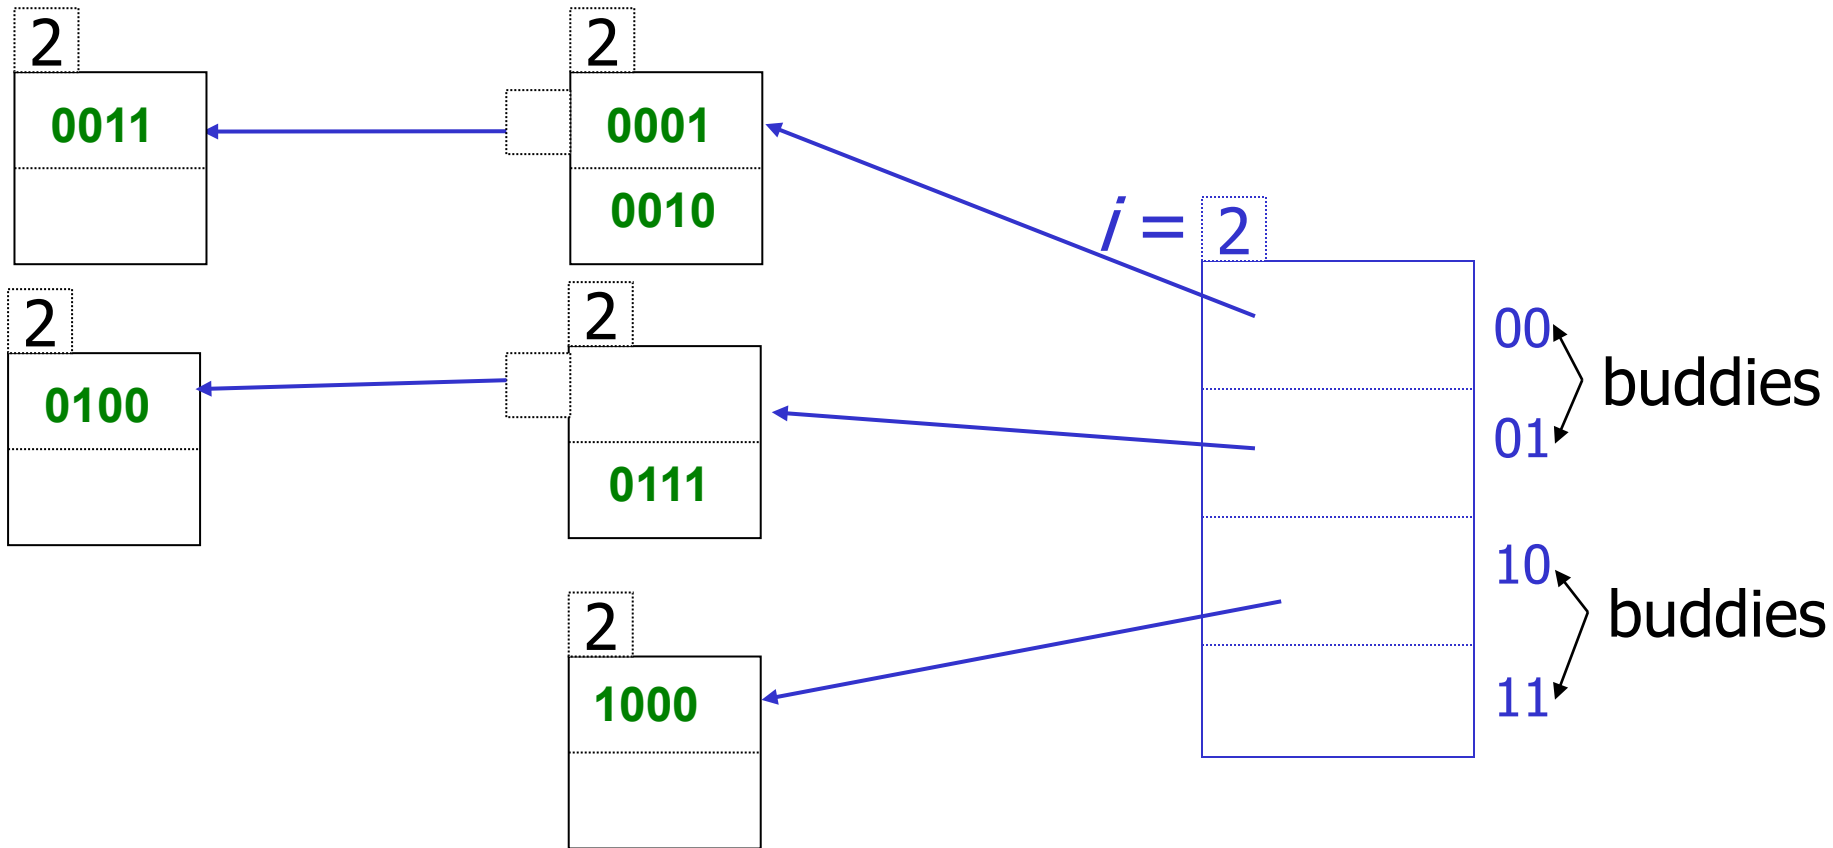
Insert 1000 = 8

Insert 0011 = 3



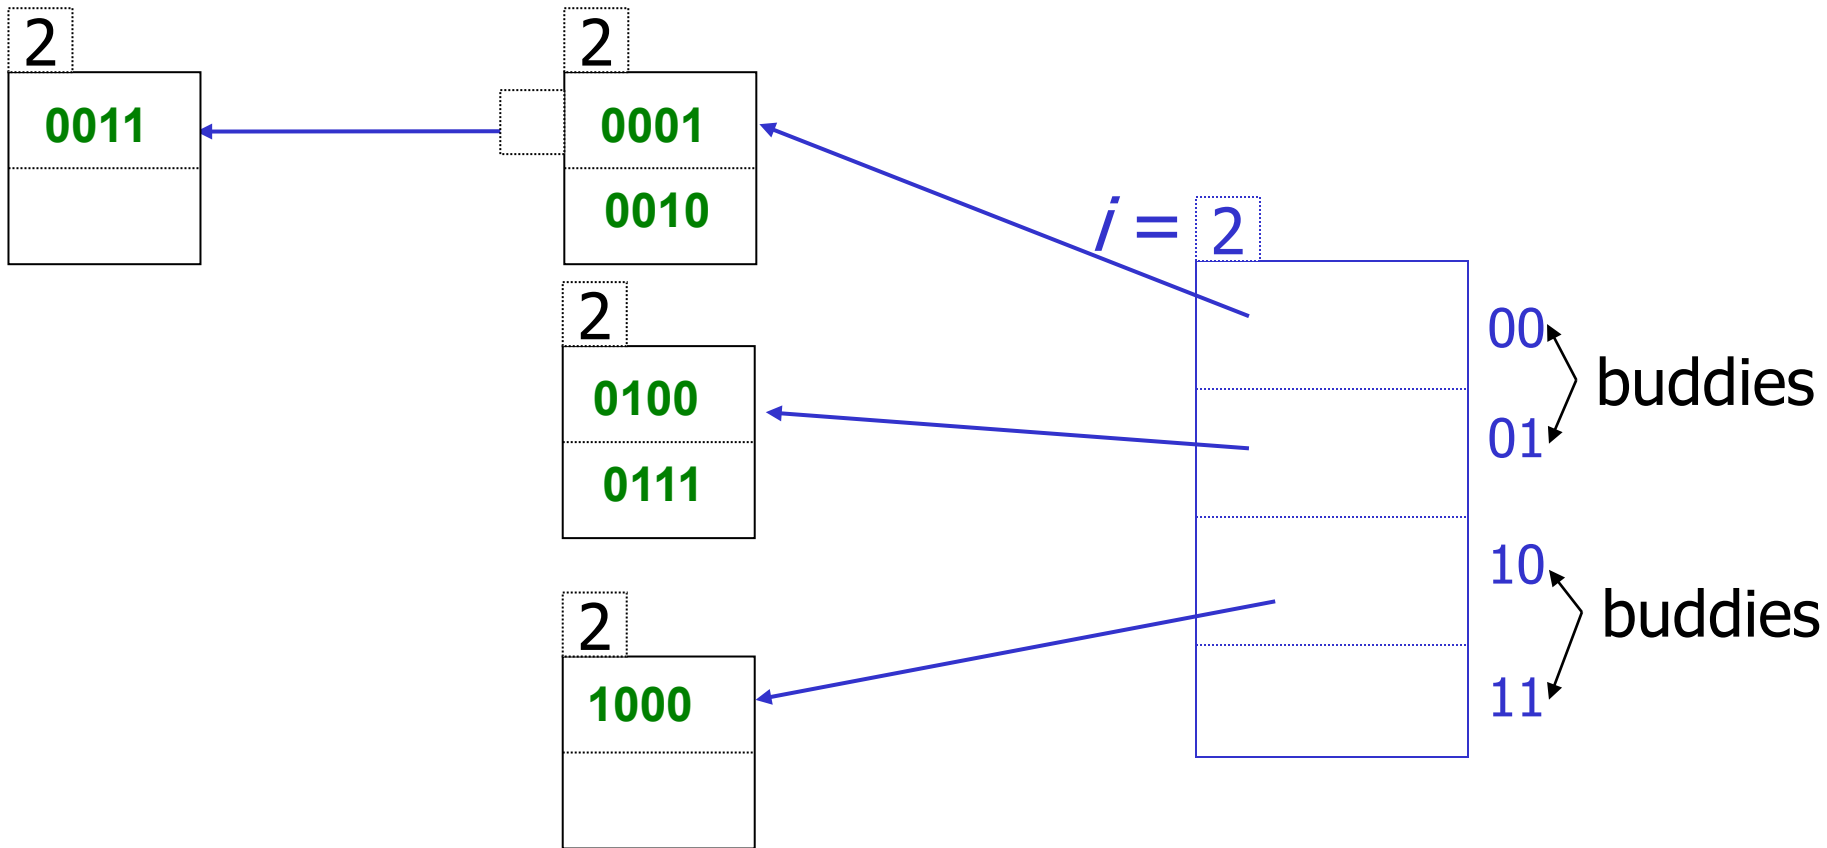
Delete **1001** = 9

Delete **0110** = 6




Delete 1001 = 9

Delete 0110 = 6



Delete 1001 = 9


Delete 0110 = 6



New directory

Delete 0100 = 4

Delete 0001 = 1



Delete **0010 = 2**

Delete **1000 = 8**

- Problem 10
 - For this problem I have set the *work_mem* = '1.5GB'
 - Appropriate index: create index worksFor_btree on worksfor USING btree(cname);

Before Any index			After B+ Tree index			After Hash Index		
	Data Output	Explain	Messages	Notifications	Scratch Pad		Data Output	Explain
	size of relation worksFor integer	Avg execution time to scan(in ms) numeric					size of relation worksFor integer	Avg execution time to scan(in ms) numeric
1	1000	41.857				1	1000	46.582
2	10000	41.348				2	10000	47.052
3	100000	50.180				3	100000	61.007
4	1000000	138.256				4	1000000	104.000

- Problem 11
 - Appropriate index: create index worksFor_salary_hash on worksfor using hash(salary);

Before Any index			After B+ Tree index			After Hash Index		
	Data Output	Explain	Messages	Notifications	Scratch Pad		Data Output	Explain
	size of relation worksFor integer	Avg execution time to scan(in ms) numeric					size of relation worksFor integer	Avg execution time to scan(in ms) numeric
1	1000	34.374				1	1000	133.368
2	10000	57.882				2	10000	142.187
3	100000	36.252				3	100000	143.561
4	1000000	78.448				4	1000000	222.658

- Problem 12
 - Appropriate index:
 - create index worksFor_hash on worksfor using hash(cname);
 - create index worksFor_salary_hash on worksfor using hash(salary);

Before Any index			After B+ Tree index			After Hash Index		
	Data Output	Explain	Messages	Notifications	Scratch Pad		Data Output	Explain
	size of relation worksFor integer	Avg execution time to scan(in ms) numeric					size of relation worksFor integer	Avg execution time to scan(in ms) numeric
1	1000	86.075				1	1000	0.028
2	10000	59.432				2	10000	0.011
3	100000	61.236				3	100000	0.034
4	1000000	126.713				4	1000000	0.038

- Problem 13
 - Appropriate index:
 - create index worksFor_hash on worksfor using hash(cname);
 - create index worksFor_pid_hash on worksfor using hash(pid);

Before Any index	After B+ Tree index	After Hash Index
------------------	---------------------	------------------

Data Output		Explain	Messages	Notifications	Scratch Pad
	size of relation worksFor integer		Avg execution time to scan(in ms) numeric		
1		1000		48.583	
2		10000		65.926	
3		100000		122.843	
4		1000000		630.405	

Data Output		Explain	Messages	Notifications	Scratch Pad
	size of relation worksFor integer		Avg execution time to scan(in ms) numeric		
1		1000		90.521	
2		10000		94.966	
3		100000		113.782	
4		1000000		364.257	

Data Output		Explain	Messages	Notifications	Scratch Pad
	size of relation worksFor integer		Avg execution time to scan(in ms) numeric		
1		1000		46.751	
2		10000		58.878	
3		100000		61.795	
4		1000000		98.926	